

Github: <https://github.com/DidiZh/cs6650>

### Performance Testing Results:

#### 1. Single Instance Baseline Test

##### Hardware:

- Processor: Apple M1 Pro
- RAM: 16 GB
- OS: macOS

##### Software:

- Java: OpenJDK 23.0.2
- RabbitMQ: Latest (Docker)
- Spring Boot: 3.4.10

##### Consumer Configuration:

- Worker Threads: 16
- Prefetch Count: 500
- Rooms: 20
- AutoScale: Disabled

##### Test Results:

```
All 500,000 messages sent!
Send time: 52s
Send rate: 9615 msg/s

Waiting for queue to drain...
[16:25:13] Remaining: 0 | ACK rate: 0.0 msg/s

✅ All messages consumed!

=====
FINAL RESULTS
=====
Total messages:      500,000
Send time:           52s (9615 msg/s)
Consume time:        0s
Total time:          52s
Overall throughput:  9615 msg/s
Completed at:        Fri Oct 31 16:25:13 PDT 2025
=====

Results saved to: results_500k_optimized.csv
```

##### Analysis:

The system achieved an impressive 9,615 msg/s throughput with zero queue accumulation, indicating that the consumer processing speed matched or exceeded the publishing rate. This demonstrates excellent system efficiency and proper configuration.

##### Key Observations:

- Real-time Processing: Consume time of 0s indicates messages were processed as fast as they arrived
- Zero Queue Buildup: Peak queue depth of 0 shows perfect balance between publish and

consume rates

- Stable Performance: No errors or channel issues throughout the entire test
- High Throughput: 9,615 msg/s represents strong performance for a single consumer instance

## 2. Load Balanced - 2 Instances Test

Instance 1: rooms 1-10, 16 threads

Instance 2: rooms 11-20, 16 threads

Total worker threads: 32

```
| Metric | Value |
|-----|-----|
| Total Messages | 500,000 |
| Send Time | 53 seconds |
| Throughput | 9,433 msg/s |
| Consume Time | 0 seconds |
| Total Time | 53 seconds |
```

Comparison with single instance:

```
| Metric | 1 Instance | 2 Instances | Change |
|-----|-----|-----|-----|
| Throughput | 9,615 msg/s | 9,433 msg/s | -1.9% |
| Total Time | 52s | 53s | +1s |
```

Analysis:

The 2-instance configuration maintained similar performance to the single instance baseline. The consume time of 0s in both tests indicates that consumer capacity exceeds the publishing rate, revealing that the bottleneck lies in the message publishing method rather than consumer processing capacity.

## 3. Load Balanced - 4 Instances Test:

- 4 instances, each handling 5 rooms
- 16 threads per instance
- Total worker threads: 64

```
| Metric | Value |
|-----|-----|
| Total Messages | 500,000 |
| Send Time | 51 seconds |
| Throughput | 9,804 msg/s |
| Consume Time | 0 seconds |
| Total Time | 51 seconds |
```

Scalability analysis:

```
| Configuration | Threads | Throughput | Scaling Factor |
|-----|-----|-----|-----|
| 1 Instance | 16 | 9,615 msg/s | 1.0x |
| 2 Instances | 32 | 9,433 msg/s | 0.98x |
| 4 Instances | 64 | 9,804 msg/s | 1.02x |
```


Key finding:

Consumer Capacity Not the Bottleneck: All three configurations achieved 0s consume time, indicating consumers processed messages as fast as they arrived

Throughput remained consistently around 9,500-9,800 msg/s across all configurations, limited by the message publishing script's capacity

## System Design Documentation:

### 1. System Architecture



#### Server (Spring Boot):

Manages WebSocket connections from clients

Publishes messages to RabbitMQ exchange

Provides internal broadcast API for consumers

Port: 8080

#### RabbitMQ (Message Queue):

Topic exchange: chat.exchange

20 persistent queues: room.1 to room.20

Message routing based on room ID

Prefetch count: 500

#### Consumer (Java Application):

Consumes from 20 queues using dedicated channels

16 worker threads for parallel processing

AckContext ensures reliable message acknowledgment

Calls Server's broadcast API via HTTP

### 2. Key Implementation Features:

#### Reliable Message Acknowledgment:

Used AckContext to bind delivery tag with originating channel

Ensures ACK/NACK is sent on the correct channel

Prevents cross-channel ACK errors

```
public void ack(AckContext ctx, boolean multiple) throws IOException {
    Channel ch = ctx.channel();
    if (ch != null && ch.isOpen()) {
        ch.basicAck(ctx.deliveryTag(), multiple);
    }
}
```

Error Handling with NACK:

ConnectException caught and handled with NACK(requeue=true)

Prevents message loss when Server is unavailable

Messages automatically requeued for retry

```
catch (ConnectException e) {  
    queue.nack(ctx, true); // requeue for retry  
}
```

Performance Optimizations:

Disabled AutoScale:

- Problem: AutoScale caused channel closures during high load
- Solution: Fixed thread pool size (16 threads)
- Result: Stable performance, no channel errors

High Prefetch Count (500):

- Problem: Low prefetch caused frequent network round-trips
- Solution: Increased prefetch to 500
- Result: ~40% throughput improvement

Optimization Impact Summary:

Aspect	Before	After	Improvement
Publish Rate	37 msg/s	9,615 msg/s	260x
Error Rate	High (channel errors)	0%	100% reduction
Stability	Unstable (AutoScale)	Stable	Qualitative
Queue Depth	Variable	0 (real-time)	Perfect balance

### 3. Test Methodology:

Message Generation:

- Total messages: 500,000
- Distribution: Round-robin across 20 rooms
- Message format: JSON with messageId, roomId, userId, username, message, timestamp

Monitoring:

- RabbitMQ Management Console for queue metrics
- Custom monitoring script tracking:
  - Queue depth (ready + unacknowledged)
  - Publish rate (msg/s)
  - ACK rate (msg/s)
- 10-second refresh interval

Metrics Collected:

- Send time: Time to publish all messages
- Consume time: Time to process all messages
- Total time: End-to-end duration
- Throughput: Messages per second
- Peak queue depth: Maximum messages in queues

#### 4. Observations and Analysis:

##### Network I/O Overhead:

- Each message requires HTTP POST from Consumer to Server
- Network latency limits overall throughput
- Impact increases with message volume

##### RabbitMQ Queue Management:

- Queue depth accumulates during high publish rates
- Memory consumption increases with queue size

##### Message Processing Overhead:

- JSON serialization/deserialization for each message
- String encoding/decoding adds latency
- Linear impact with message volume

#### **Performance Bonus:**

Best Throughput: 9,615 msg/s (single instance)