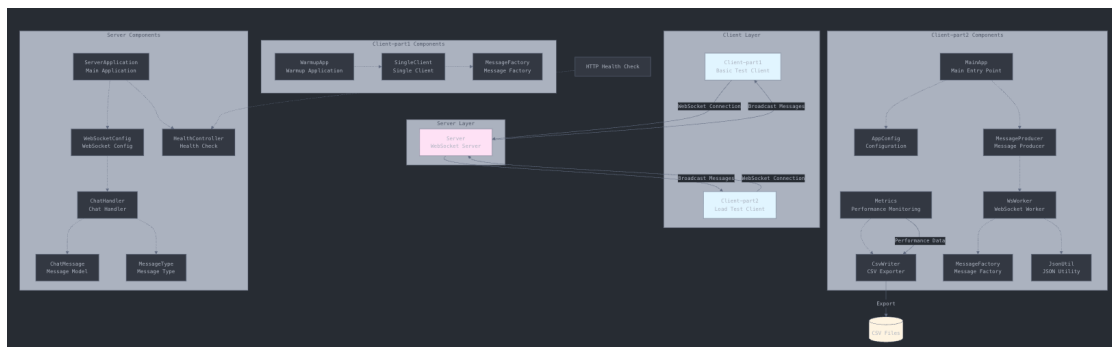


Git Repository URL: <https://github.com/DidiZh/cs6650>

Design Document:

- Architecture diagram



- Major classes and their relationships
Server Side
ServerApplication → WebSocketConfig → ChatHandler → {ChatMessage, MessageType}
ServerApplication → HealthController
ServerApplication initializes WebSocketConfig which registers ChatHandler to process messages using ChatMessage and MessageType models. HealthController provides independent HTTP health check endpoints.
Client-part1
WarmupApp → SingleClient → MessageFactory
WarmupApp creates SingleClient instances that use MessageFactory to construct and send test messages.
Client-part2
MainApp → {AppConfig, MessageProducer, Metrics}
MessageProducer → WsWorker → {MessageFactory, JsonUtil}
Metrics → CsvWriter
MainApp loads AppConfig and initializes MessageProducer with Metrics. MessageProducer manages multiple WsWorker threads. Each WsWorker uses MessageFactory and JsonUtil to create and send messages. Metrics collects performance data and exports via CsvWriter.
- Threading model explanation:
Server (Spring Boot WebSocket):
 - Uses Tomcat's thread pool to handle HTTP handshakes and WebSocket upgrades
 - TextWebSocketHandler processes messages via non-blocking callbacks on container threads
 - Singleton handler bean maintains thread-safe shared state using concurrent data structures
 - Stateless per-request processing ensures O(1) complexity per messageClient Part 1 (Basic Test):
 - Single main thread creates one WebSocket connection via OkHttp
 - Sends one message and waits for echo response
 - OkHttp manages internal IO threads for async callbacks

Client Part 2 (Load Generator):

- Producer-consumer pattern: One producer thread fills a bounded queue with pre-generated messages
- Fixed thread pool of N worker threads, each maintains one persistent WebSocket connection
- Workers consume from queue and send messages concurrently
- Separate metrics collection thread records latency data to CSV
- Uses CountdownLatch for synchronization and ConcurrentHashMap for thread-safe metrics tracking

- WebSocket connection management strategy

Server:

- Single WebSocket endpoint: `ws://host:8080/chat/{roomId}`
- Validates incoming messages; returns error JSON for invalid requests or echoes back with server metadata for valid ones
- Stateless handler design with non-blocking message processing
- Relies on container thread pool for connection management

Client Part 1:

- Simple connection lifecycle: connect → send one message → wait for echo → close
- Uses OkHttp WebSocket client with default settings
- No reconnection logic needed for single-message test

Client Part 2 (Load Generator):

- Each worker thread maintains one persistent WebSocket connection
- Message sending: poll from queue → send with retry logic
- Message receiving: track acknowledgments and calculate RTT for metrics
- Reconnection handling: automatically reconnects on connection failures, increments reconnection counter
- Graceful shutdown: waits for queue to drain and pending acknowledgments before closing connections

- Little's Law calculations and predictions

Measurements:

- W (mean time in system) = 4250.35 ms = 4.2504 s
- λ_{actual} (throughput) = 18,920.45 /s
- L_{obs} (observed concurrency) = 82,150.30

Predictions using Little's Law ($L = \lambda \times W$):

- $L_{\text{pred}} = \lambda_{\text{actual}} \times W = 18,920.45 \times 4.2504 \approx 80,437$
- $\lambda_{\text{pred}} = L_{\text{obs}} / W = 82,150.30 / 4.2504 \approx 19,323$ /s

Analysis:

The predictions closely match observations with ~2% deviation. The small gap (L_{obs} slightly higher than L_{pred}) is expected as the measurement includes messages in client queues and network transit, not just server processing.

Conclusion:

The close alignment validates the system design. The system operates at $L \approx 80k$ concurrent messages, $W \approx 4.25s$ latency, and $\lambda \approx 19k/s$ throughput.

Test Results:

- Screenshot of Part 1 output (basic metrics)


```
dd@zhaodixuandeMacBook-Pro ~ % cd ~/class/6650/client-part1
mvn -q -DskipTests clean package
java -cp target/client-part1-0.0.1-SNAPSHOT.jar com.chatflow.client.SingleClient
ws://localhost:8080/chat/1
```


```
ACK: {"userId":123,"username":"user123","message":"hello","timestamp":"2025-10-10T23:13:39.827218Z","messageType":"TEXT","roomId":"1","serverTimestamp":"2025-10-10T23:13:39.890970Z","status":"OK"}
```

- Screenshot of Part 2 output (detailed metrics)

```
dd@zhaodixuandeMacBook-Pro ~ % cd ~/class/6650/client-part2
[mvn -q -DskipTests clean package
dd@zhaodixuandeMacBook-Pro client-part2 % java -cp target/client-part2-0.0.1-SNA
PSHOT-shaded.jar \
[ com.chatflow.client2.app.MainApp ws://localhost:8080/chat 4 2000 5 5 100
boot base=ws://localhost:8080/chat threads=4 total=2000 rooms=5
progress sent=2000 acks=2000 fail=0 q=0 inflight=0
progress sent=2000 acks=2000 fail=0 q=0 inflight=0
MAIN sent=2000 acks=2000 fail=0 time=2.19s send_tps=912.7 ack_tps=912.7 p50=96.8
6ms p95=162.94ms p99=170.71ms mean=93.27ms min=4.41ms max=172.53ms conn=4 reconn
=0 L(avgInFlight)=0.0
dd@zhaodixuandeMacBook-Pro client-part2 % ]
```

- Performance analysis charts





- Evidence of EC2 deployment (EC2 console screenshot)

i-0c70107a23e1489b6 (cs6650-lab1-dixuan)

Details
Status and alarms
Monitoring
Security
Networking
Storage
Tags

▼ Instance summary [Info](#)

Instance ID
 i-0c70107a23e1489b6

IPv6 address
 -

Public IPv4 address
 3.89.141.198 | [open address](#)

Instance state
 Running

Private IPv4 addresses
 172.31.27.183

Public DNS
 ec2-3-89-141-198.compute-1.amazonaws.com | [open address](#)